

Making predictions

INTRODUCTION TO REGRESSION IN R



Richie Cotton

Learning Solutions Architect at
DataCamp

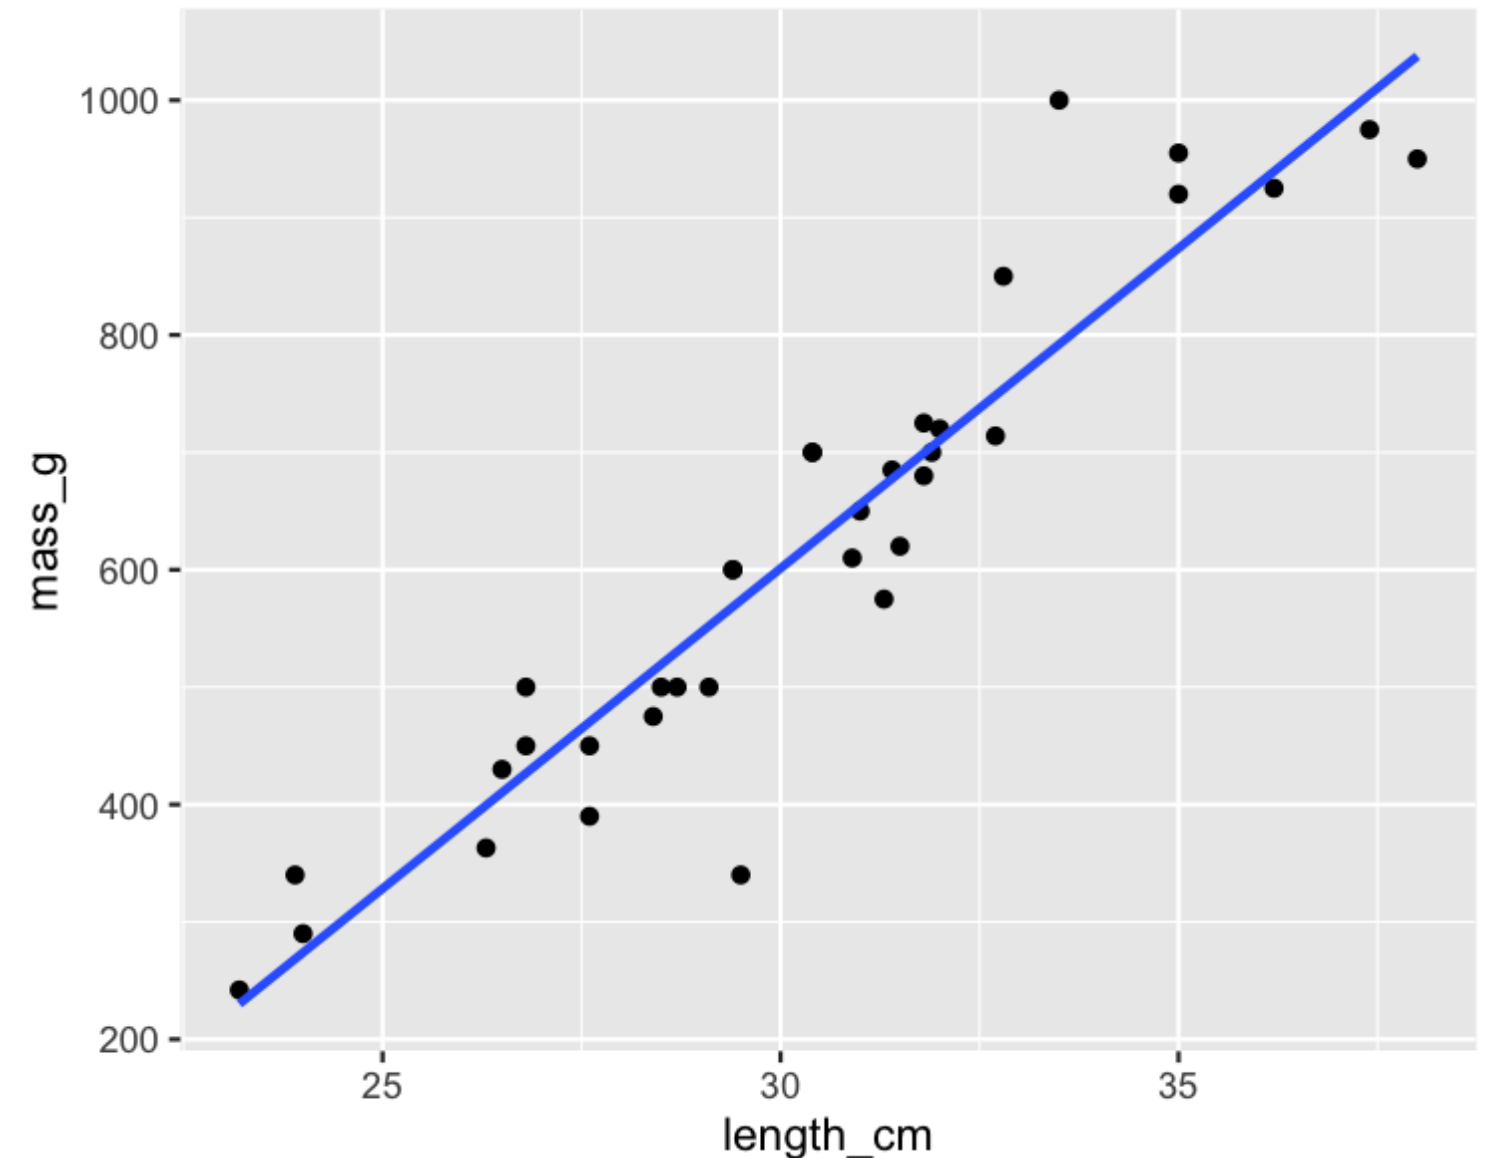
The fish dataset: bream

```
bream <- fish %>%  
  filter(species == "Bream")
```

species	length_cm	mass_g
Bream	23.2	242
Bream	24.0	290
Bream	23.9	340
Bream	26.3	363
Bream	26.5	430
...

Plotting mass vs. length

```
ggplot(bream, aes(length_cm, mass_g)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE)
```



Running the model

```
mdl_mass_vs_length <- lm(mass_g ~ length_cm, data = bream)
```

Call:

```
lm(formula = mass_g ~ length_cm, data = bream)
```

Coefficients:

(Intercept)	length_cm
-1035.35	54.55

Data on explanatory values to predict

If I set the explanatory variables to these values,
what value would the response variable have?

```
library(dplyr)
explanatory_data <- tibble(length_cm = 20:40)
```

Call predict()

```
library(tibble)
explanatory_data <- tibble(length_cm = 20:40)
```

```
predict mdl_mass_vs_length, explanatory_data)
```

1	2	3	4	5	6
55.65205	110.20203	164.75202	219.30200	273.85198	328.40196
7	8	9	10	11	12
382.95194	437.50192	492.05190	546.60188	601.15186	655.70184
13	14	15	16	17	18
710.25182	764.80181	819.35179	873.90177	928.45175	983.00173
19	20	21			
1037.55171	1092.10169	1146.65167			

Predicting inside a data frame

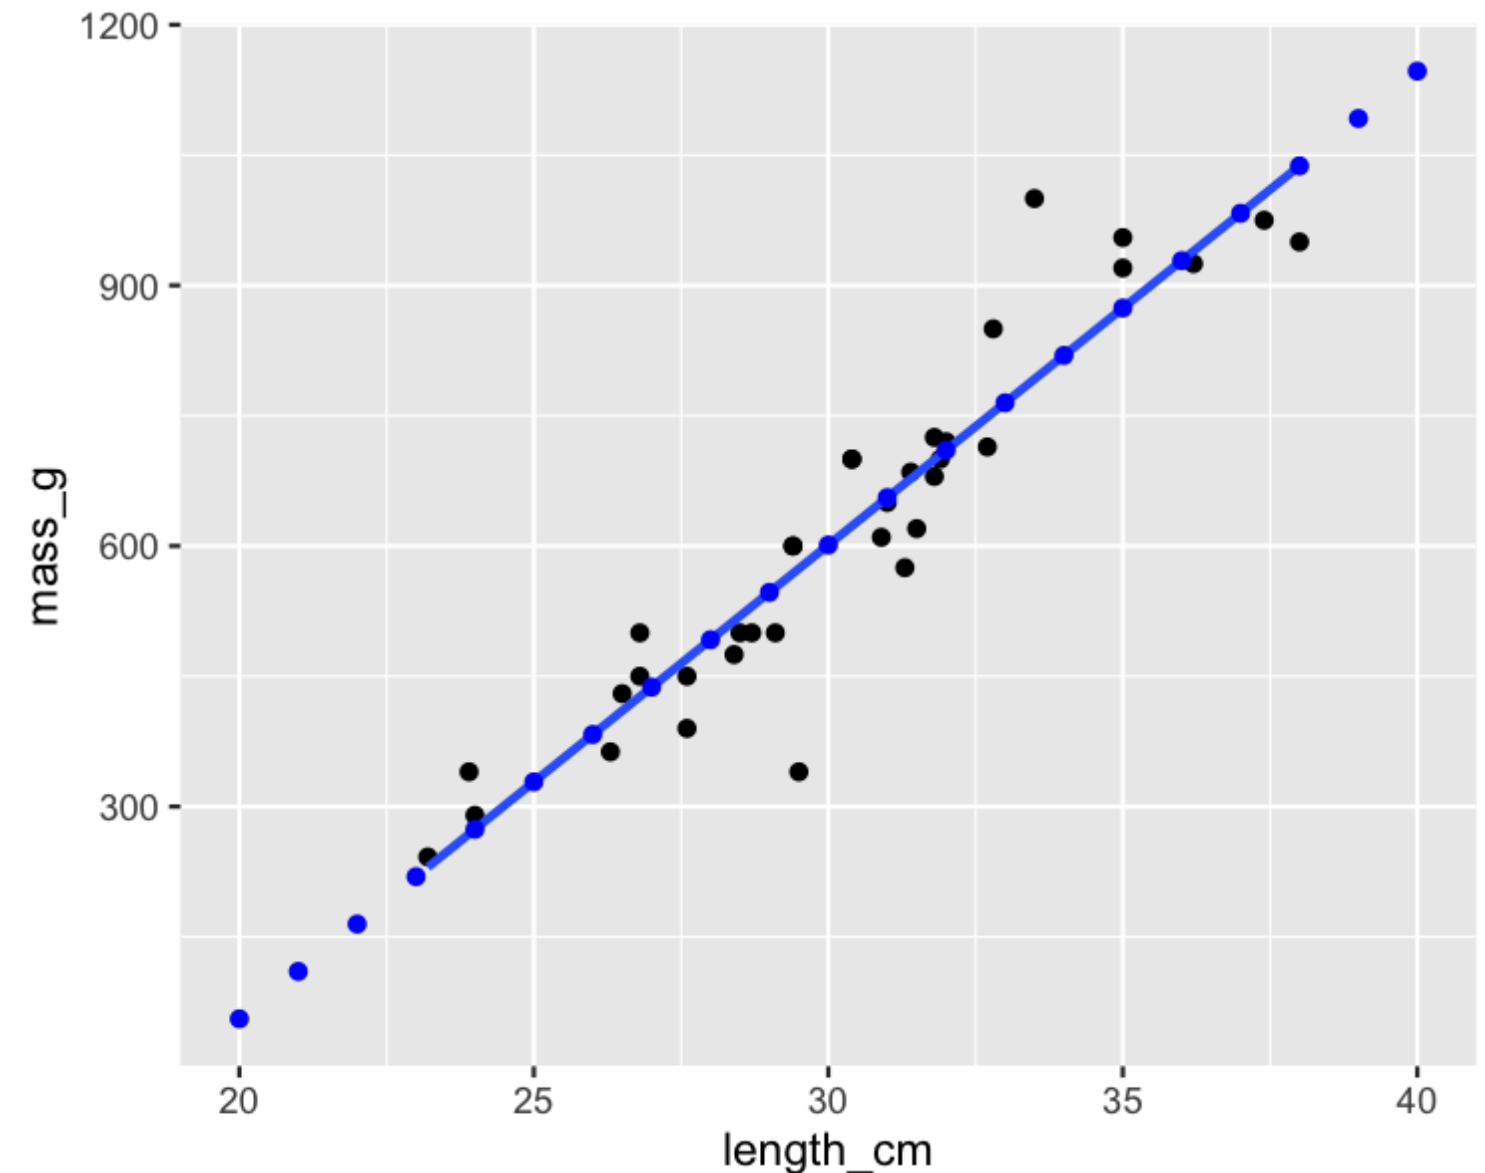
```
library(dplyr)
explanatory_data <- tibble(length_cm = 20:40)
```

```
prediction_data <- explanatory_data %>%
  mutate(
    mass_g = predict(
      mdl_mass_vs_length, explanatory_data
    )
  )
```

```
# A tibble: 21 x 2
  length_cm mass_g
  <int>    <dbl>
1      20    55.7
2      21   110.
3      22   165.
4      23   219.
5      24   274.
6      25   328.
7      26   383.
8      27   438.
9      28   492.
10     29   547.
# ... with 11 more rows
```

Showing predictions

```
ggplot(bream, aes(length_cm, mass_g)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE) +  
  geom_point(  
    data = prediction_data,  
    color = "blue"  
  )
```

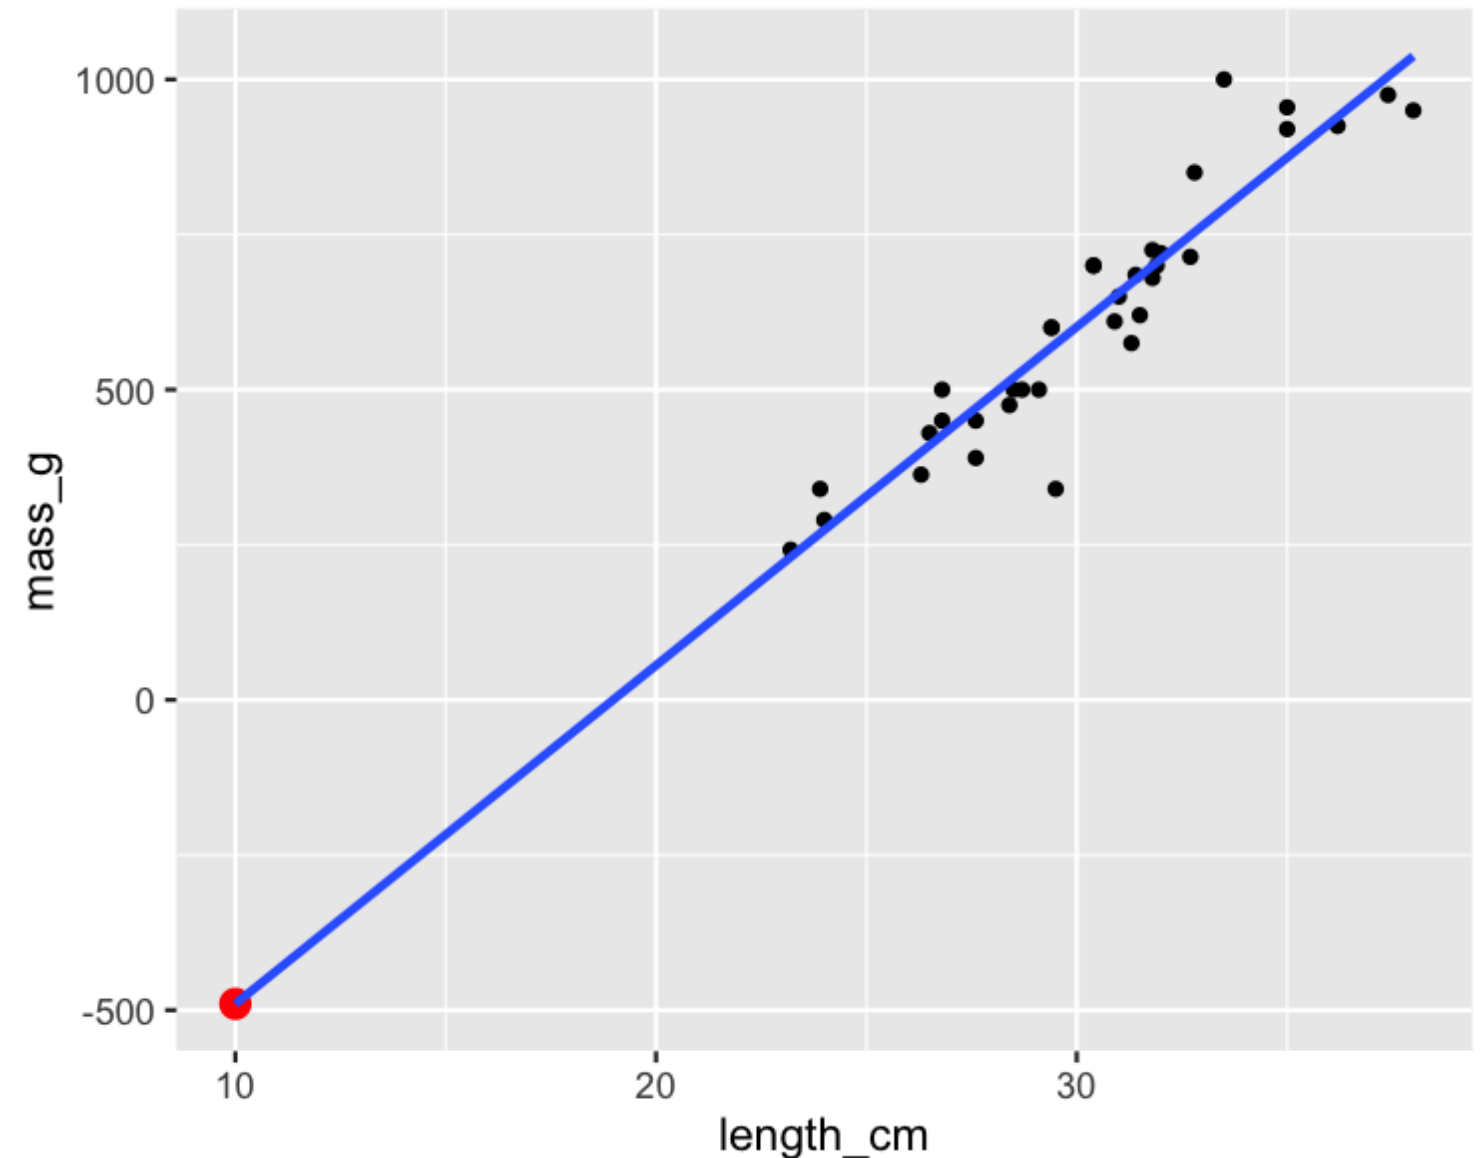


Extrapolating

Extrapolating means making predictions outside the range of observed data.

```
explanatory_little_bream <- tibble(length_cm = 10)
explanatory_little_bream %>%
  mutate(
    mass_g = predict(
      mdl_mass_vs_length, explanatory_little_bream
    )
  )
```

```
# A tibble: 1 x 2
  length_cm mass_g
  <dbl>    <dbl>
1      10 -490.
```



Let's practice!

INTRODUCTION TO REGRESSION IN R

Working with model objects

INTRODUCTION TO REGRESSION IN R



Richie Cotton

Learning Solutions Architect at
DataCamp

coefficients()

```
mdl_mass_vs_length <- lm(mass_g ~ length_cm, data = bream)
```

```
Call:
lm(formula = mass_g ~ length_cm, data = bream)

Coefficients:
(Intercept)    length_cm
   -1035.35         54.55
```

```
coefficients(mdl_mass_vs_length)
```

```
(Intercept)    length_cm
-1035.34757     54.54998
```

fitted()

fitted values: predictions on the original dataset

```
fitted(mdl_mass_vs_length)
```

or equivalently

```
explanatory_data <- bream %>%  
  select(length_cm)  
  
predict(mdl_mass_vs_length, explanatory_data)
```

1	2	3	4	5
230.2120	273.8520	268.3970	399.3169	410.2269
6	7	8	9	10
426.5919	426.5919	470.2319	470.2319	519.3269
11	12	13	14	15
513.8719	530.2369	552.0569	573.8769	568.4219
16	17	18	19	20
568.4219	622.9719	622.9719	650.2468	655.7018
21	22	23	24	25
672.0668	677.5218	682.9768	699.3418	704.7968
26	27	28	29	30
699.3418	710.2518	748.4368	753.8918	792.0768
31	32	33	34	35
873.9018	873.9018	939.3617	1004.8217	1037.5517

residuals()

Residuals: actual response values minus predicted response values

```
residuals mdl_mass_vs_length
```

or equivalently

```
bream$mass_g - fitted mdl_mass_vs_length
```

1	2	3	4	5
11.788	16.148	71.603	-36.317	19.773
6	7	8	9	10
23.408	73.408	-80.232	-20.232	-19.327
11	12	13	14	15
-38.872	-30.237	-52.057	-233.877	31.578
16	17	18	19	20
31.578	77.028	77.028	-40.247	-5.702
21	22	23	24	25
-97.067	7.478	-62.977	-19.342	-4.797
26	27	28	29	30
25.658	9.748	-34.437	96.108	207.923
31	32	33	34	35
46.098	81.098	-14.362	-29.822	-87.552

summary()

```
summary mdl_mass_vs_length
```

```
Call:
lm(formula = mass_g ~ length_cm, data = bream)

Residuals:
    Min       1Q   Median       3Q      Max
-233.9   -35.4    -4.8    31.6   207.9

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -1035.35     107.97   -9.59  4.6e-11 ***
length_cm     54.55       3.54   15.42 < 2e-16 ***

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 74.2 on 33 degrees of freedom
Multiple R-squared:  0.878,    Adjusted R-squared:  0.874
F-statistic: 238 on 1 and 33 DF,  p-value: <2e-16
```

summary(): call

```
Call:  
lm(formula = mass_g ~ length_cm, data = bream)
```


summary(): residuals

Residuals:

Min	1Q	Median	3Q	Max
-233.9	-35.4	-4.8	31.6	207.9

summary(): coefficients

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1035.35	107.97	-9.59	4.6e-11	***
length_cm	54.55	3.54	15.42	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

summary(): model metrics

```
Residual standard error: 74.2 on 33 degrees of freedom  
Multiple R-squared:  0.878,    Adjusted R-squared:  0.874  
F-statistic: 238 on 1 and 33 DF,  p-value: <2e-16
```

tidy()

```
library(broom)
```

```
tidy mdl_mass_vs_length
```

```
# A tibble: 2 x 5
  term          estimate std.error statistic  p.value
<chr>          <dbl>     <dbl>     <dbl>    <dbl>
1 (Intercept) -1035.      108.      -9.59 4.58e-11
2 length_cm     54.5       3.54      15.4 1.22e-16
```

augment()

```
augment(mdl_mass_vs_length)
```

```
# A tibble: 35 x 9
  mass_g length_cm .fitted .se.fit .resid   .hat .sigma .cooksd .std.resid
  <dbl>   <dbl>   <dbl>   <dbl> <dbl>   <dbl> <dbl>   <dbl>   <dbl>
1    242     23.2    230.    28.1   11.8 0.144   75.3 0.00247    0.172
2    290     24     274.    25.6   16.1 0.119   75.2 0.00364    0.232
3    340     23.9    268.    25.9   71.6 0.122   74.1 0.0738     1.03
4    363     26.3    399.    18.9  -36.3 0.0651  75.0 0.00894   -0.507
5    430     26.5    410.    18.4   19.8 0.0616  75.2 0.00248    0.275
6    450     26.8    427.    17.6   23.4 0.0566  75.2 0.00317    0.325
7    500     26.8    427.    17.6   73.4 0.0566  74.1 0.0311     1.02
8    390     27.6    470.    15.8  -80.2 0.0452  73.9 0.0291    -1.11
9    450     27.6    470.    15.8  -20.2 0.0452  75.2 0.00185   -0.279
10   500     28.5    519.    14.1  -19.3 0.0360  75.2 0.00132   -0.265
# ... with 25 more rows
```

glance()

```
glance mdl_mass_vs_length)
```

```
# A tibble: 1 x 11
  r.squared adj.r.squared sigma statistic p.value    df logLik   AIC    BIC deviance df.residual
  <dbl>      <dbl> <dbl>      <dbl>   <dbl> <int> <dbl> <dbl> <dbl>    <dbl>      <int>
1  0.878      0.874  74.2       238. 1.22e-16     2  -199.  405.  409.  181452.     33
```

Let's practice!

INTRODUCTION TO REGRESSION IN R

Regression to the mean

INTRODUCTION TO REGRESSION IN R



Richie Cotton

Learning Solutions Architect

The concept

- Response value = fitted value + residual
- "The stuff you explained" + "the stuff you couldn't explain"
- Residuals exist due to problems in the model *and* fundamental randomness
- Extreme cases are often due to randomness
- *Regression to the mean* means extreme cases don't persist over time

Pearson's father son dataset

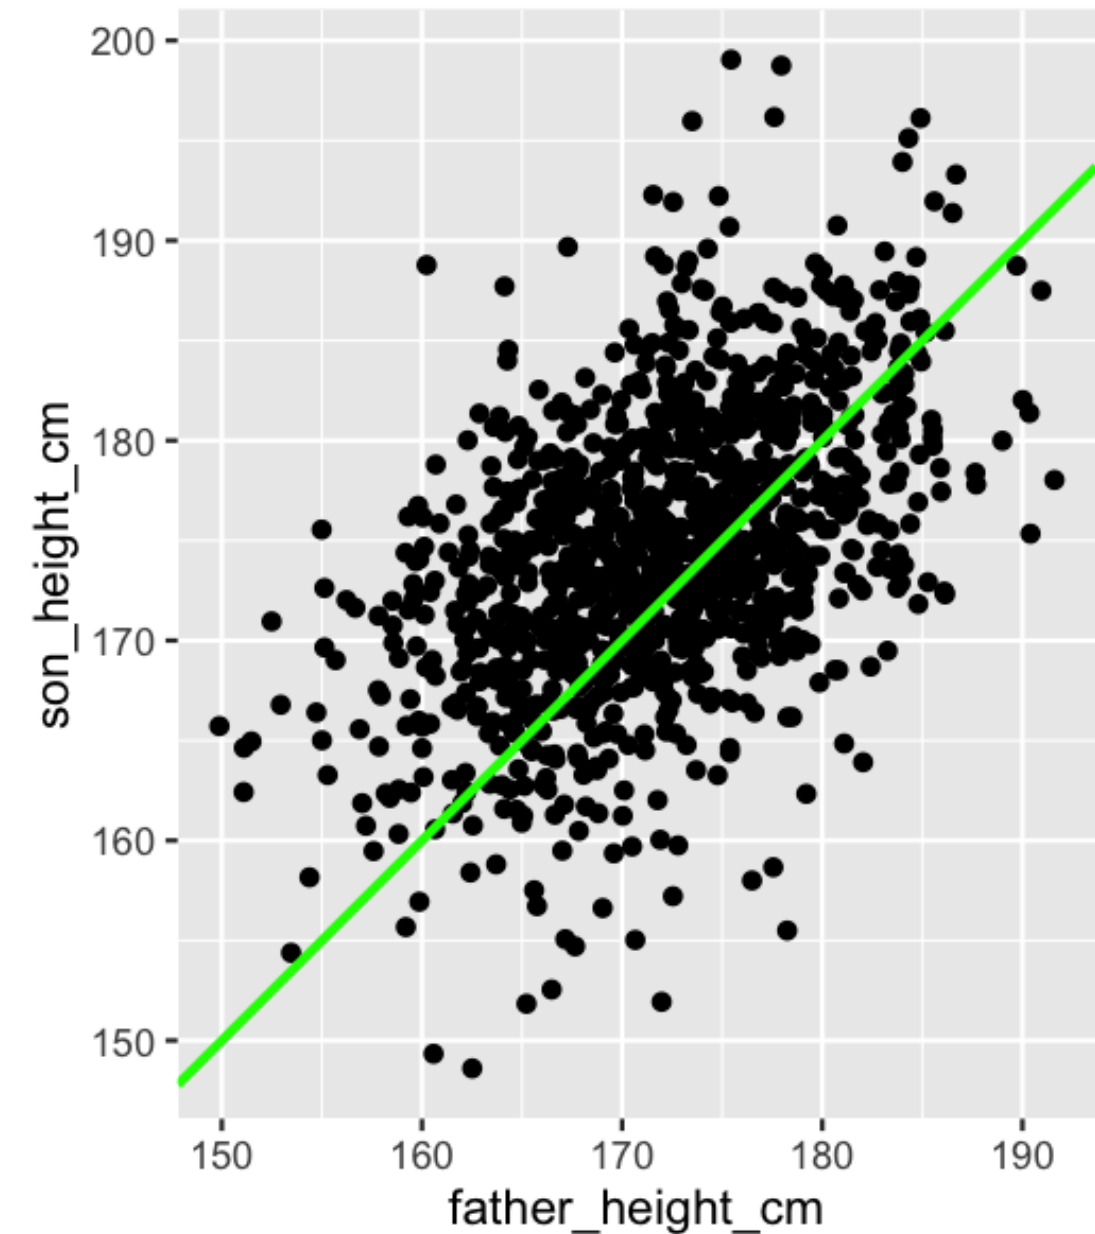
- 1078 father/son pairs
- Do tall fathers have tall sons?

father_height_cm	son_height_cm
165.2	151.8
160.7	160.6
165.0	160.9
167.0	159.5
155.3	163.3
...	...

¹ Adapted from <https://www.rdocumentation.org/packages/UsingR/topics/father.son>

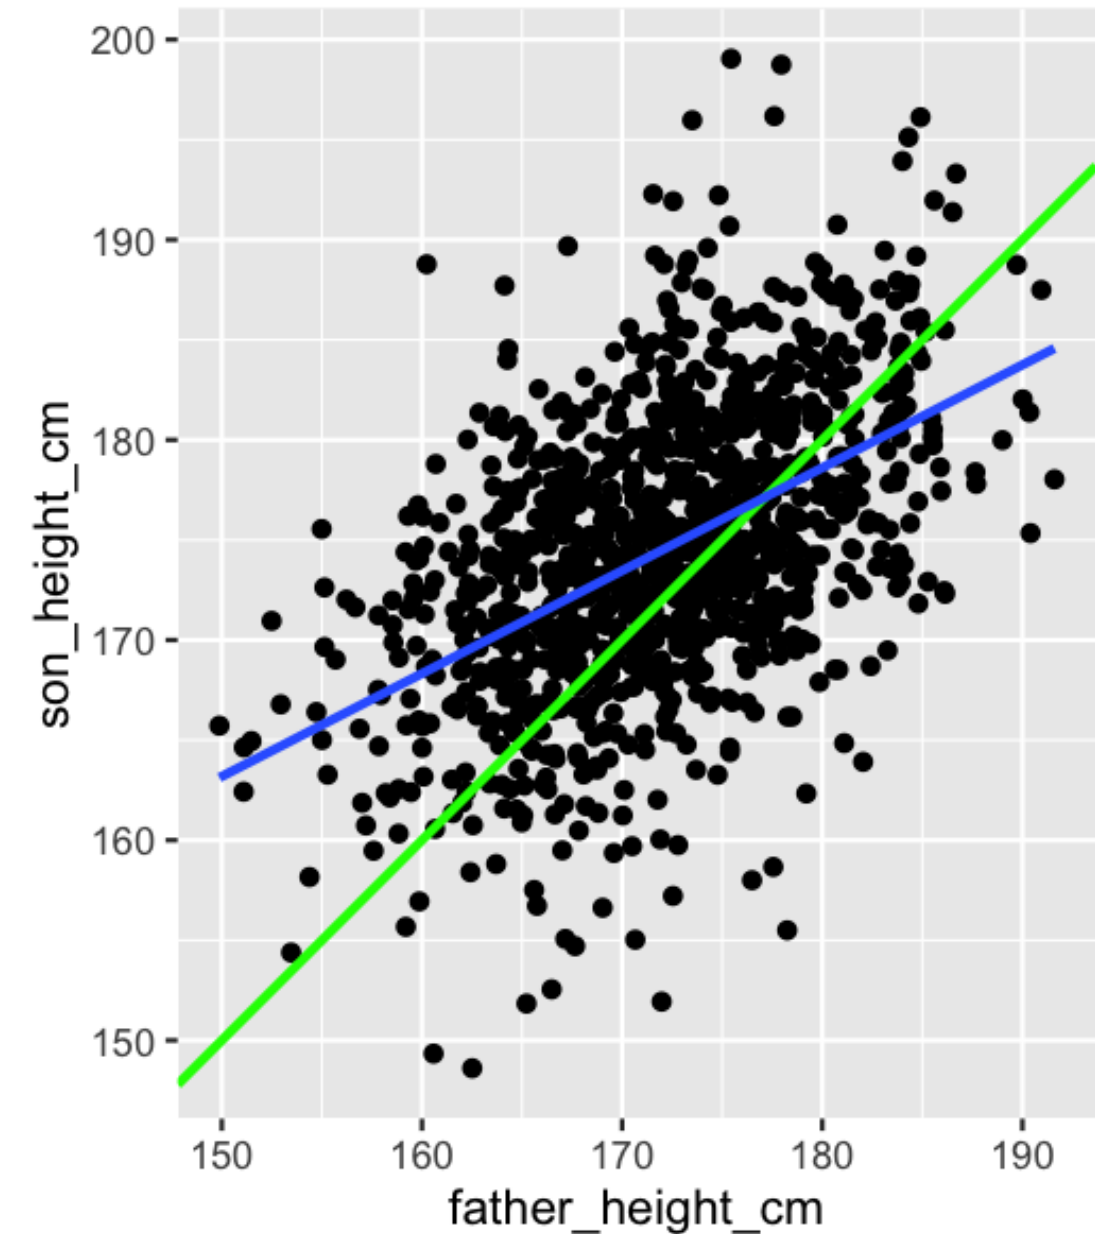
Scatter plot

```
plt_son_vs_father <- ggplot(  
  father_son,  
  aes(father_height_cm, son_height_cm)  
) +  
  geom_point() +  
  geom_abline(color = "green", size = 1) +  
  coord_fixed()
```



Adding a regression line

```
plt_son_vs_father +  
  geom_smooth(method = "lm", se = FALSE)
```



Running a regression

```
mdl_son_vs_father <- lm(  
  son_height_cm ~ father_height_cm,  
  data = father_son  
)
```

Call:

```
lm(formula = son_height_cm ~ father_height_cm, data = father_son)
```

Coefficients:

(Intercept)	father_height_cm
86.072	0.514

Making predictions

```
really_tall_father <- tibble(  
  father_height_cm = 190  
)  
predict mdl_son_vs_father, really_tall_father)
```

183.7

```
really_short_father <- tibble(  
  father_height_cm = 150  
)  
predict mdl_son_vs_father, really_short_father)
```

163.2

Let's practice!

INTRODUCTION TO REGRESSION IN R

Transforming variables

INTRODUCTION TO REGRESSION IN R



Richie Cotton

Learning Solutions Architect at
DataCamp

Perch dataset

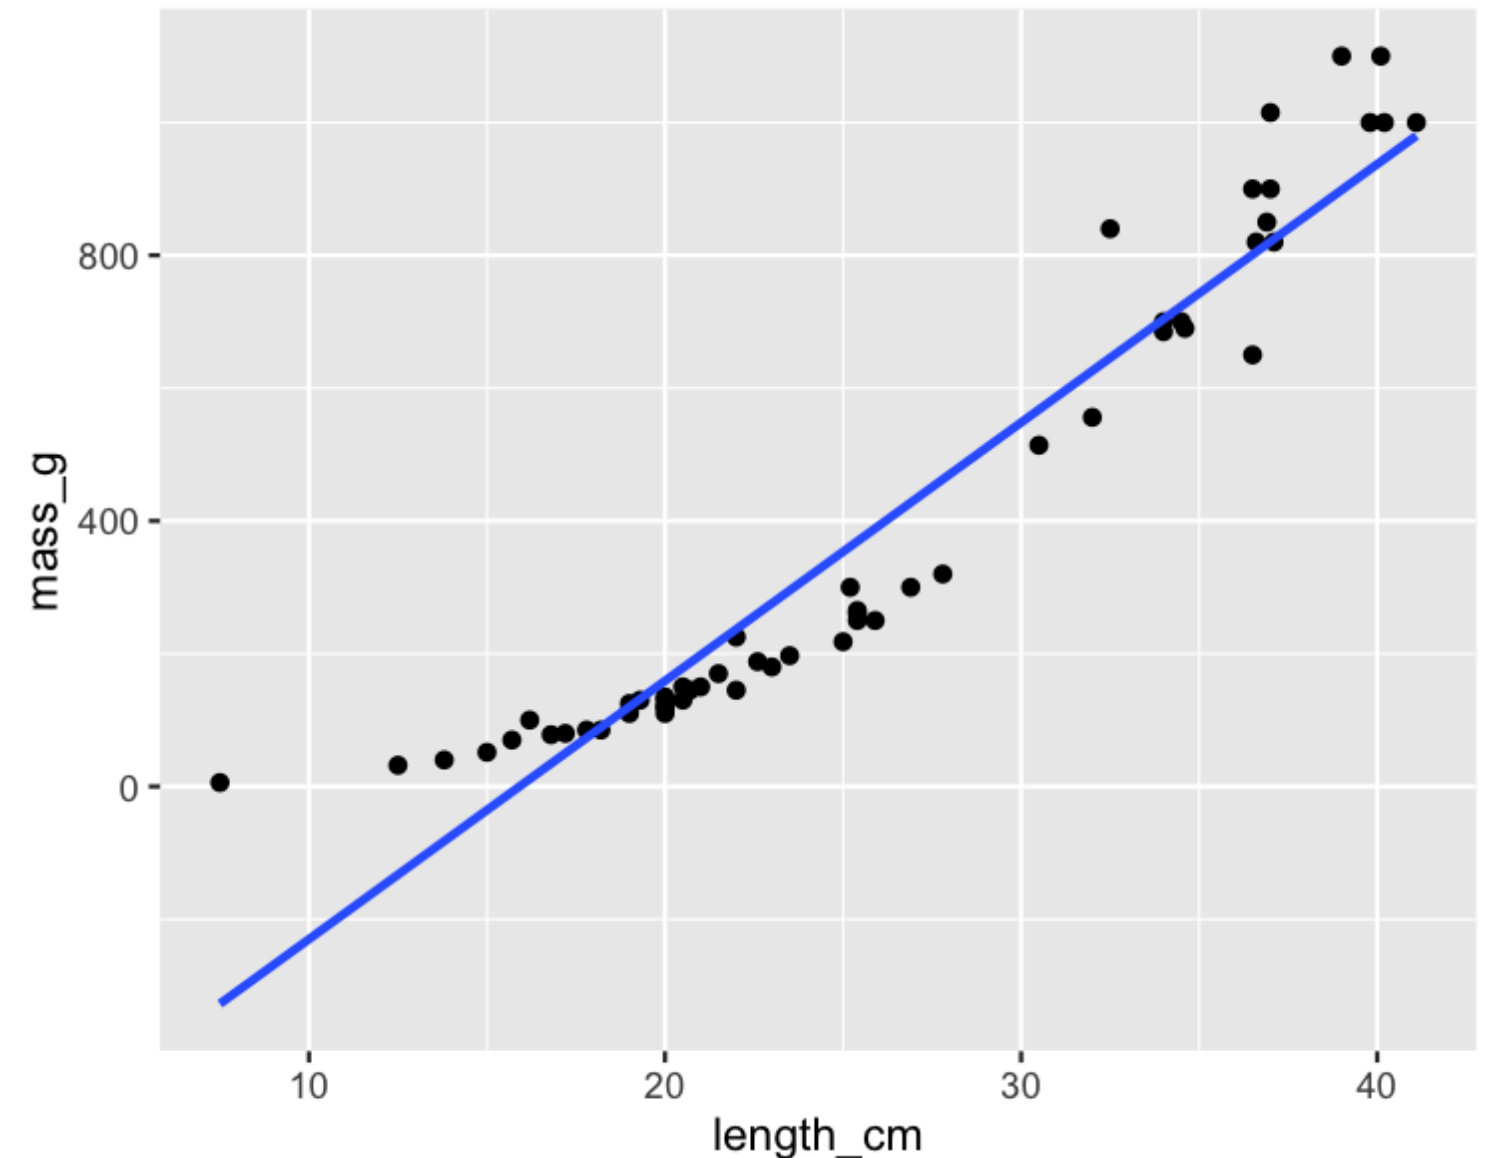
```
library(dplyr)

perch <- fish %>%
  filter(species == "Perch")
```

species	mass_g	length_cm
Perch	5.9	7.5
Perch	32.0	12.5
Perch	40.0	13.8
Perch	51.5	15.0
Perch	70.0	15.7
...

It's not a linear relationship

```
ggplot(perch, aes(length_cm, mass_g)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE)
```




Bream vs. perch



Plotting mass vs. length cubed

```
ggplot(perch, aes(length_cm ^ 3, mass_g)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE)
```



Modeling mass vs. length cubed

```
mdl_perch <- lm(mass_g ~ I(length_cm ^ 3), data = perch)
```

Call:

```
lm(formula = mass_g ~ I(length_cm^3), data = perch)
```

Coefficients:

(Intercept)	I(length_cm^3)
-0.1175	0.0168

Predicting mass vs. length cubed

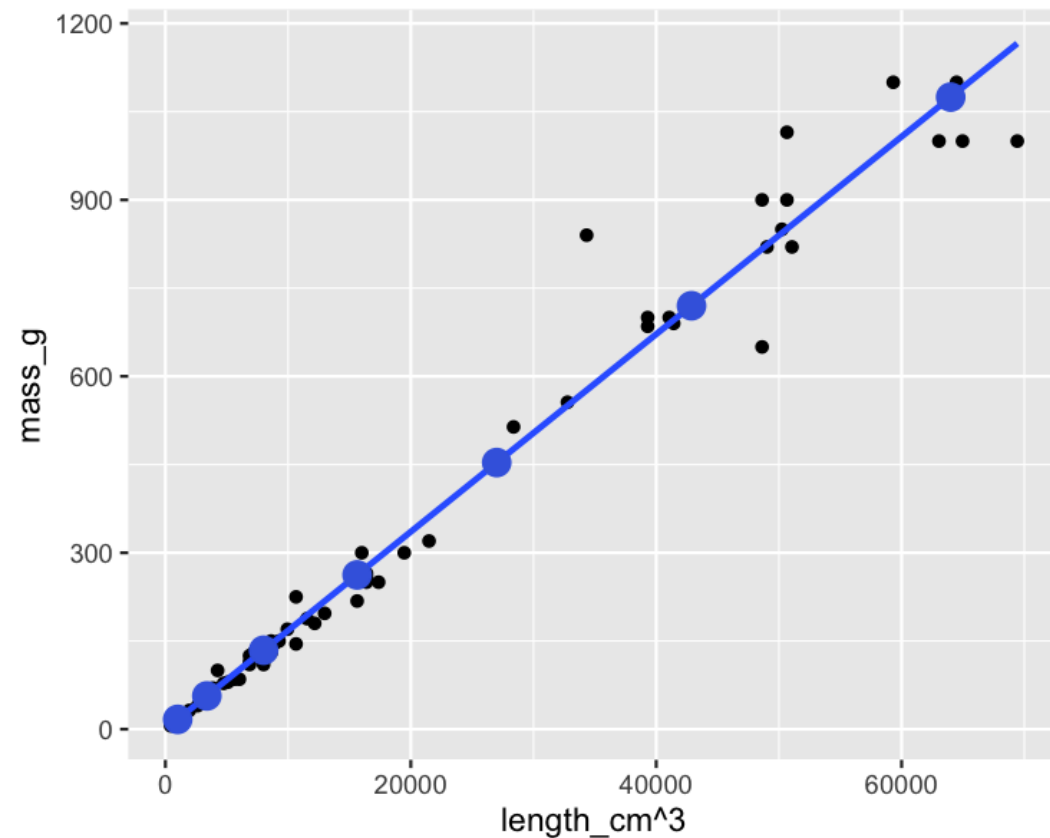
```
explanatory_data <- tibble(  
  length_cm = seq(10, 40, 5)  
)
```

```
prediction_data <- explanatory_data %>%  
  mutate(  
    mass_g = predict mdl_perch, explanatory_data)  
)
```

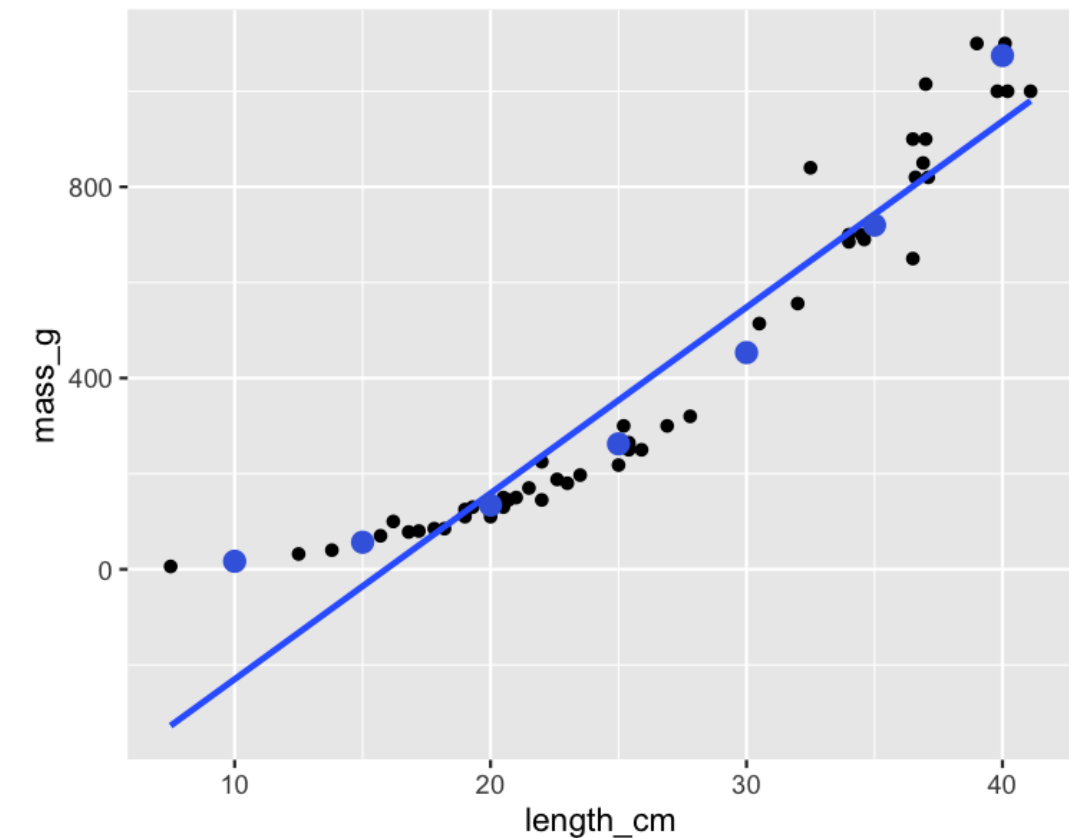
```
# A tibble: 7 x 2  
  length_cm mass_g  
    <dbl>    <dbl>  
1      10    16.7  
2      15    56.6  
3      20   134.  
4      25   262.  
5      30   453.  
6      35   720.  
7      40  1075.
```

Plotting mass vs. length cubed

```
ggplot(perch, aes(length_cm ^ 3, mass_g)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE) +  
  geom_point(data = prediction_data, color = "blue")
```



```
ggplot(perch, aes(length_cm, mass_g)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE) +  
  geom_point(data = prediction_data, color = "blue")
```



Facebook advertising dataset

How advertising works

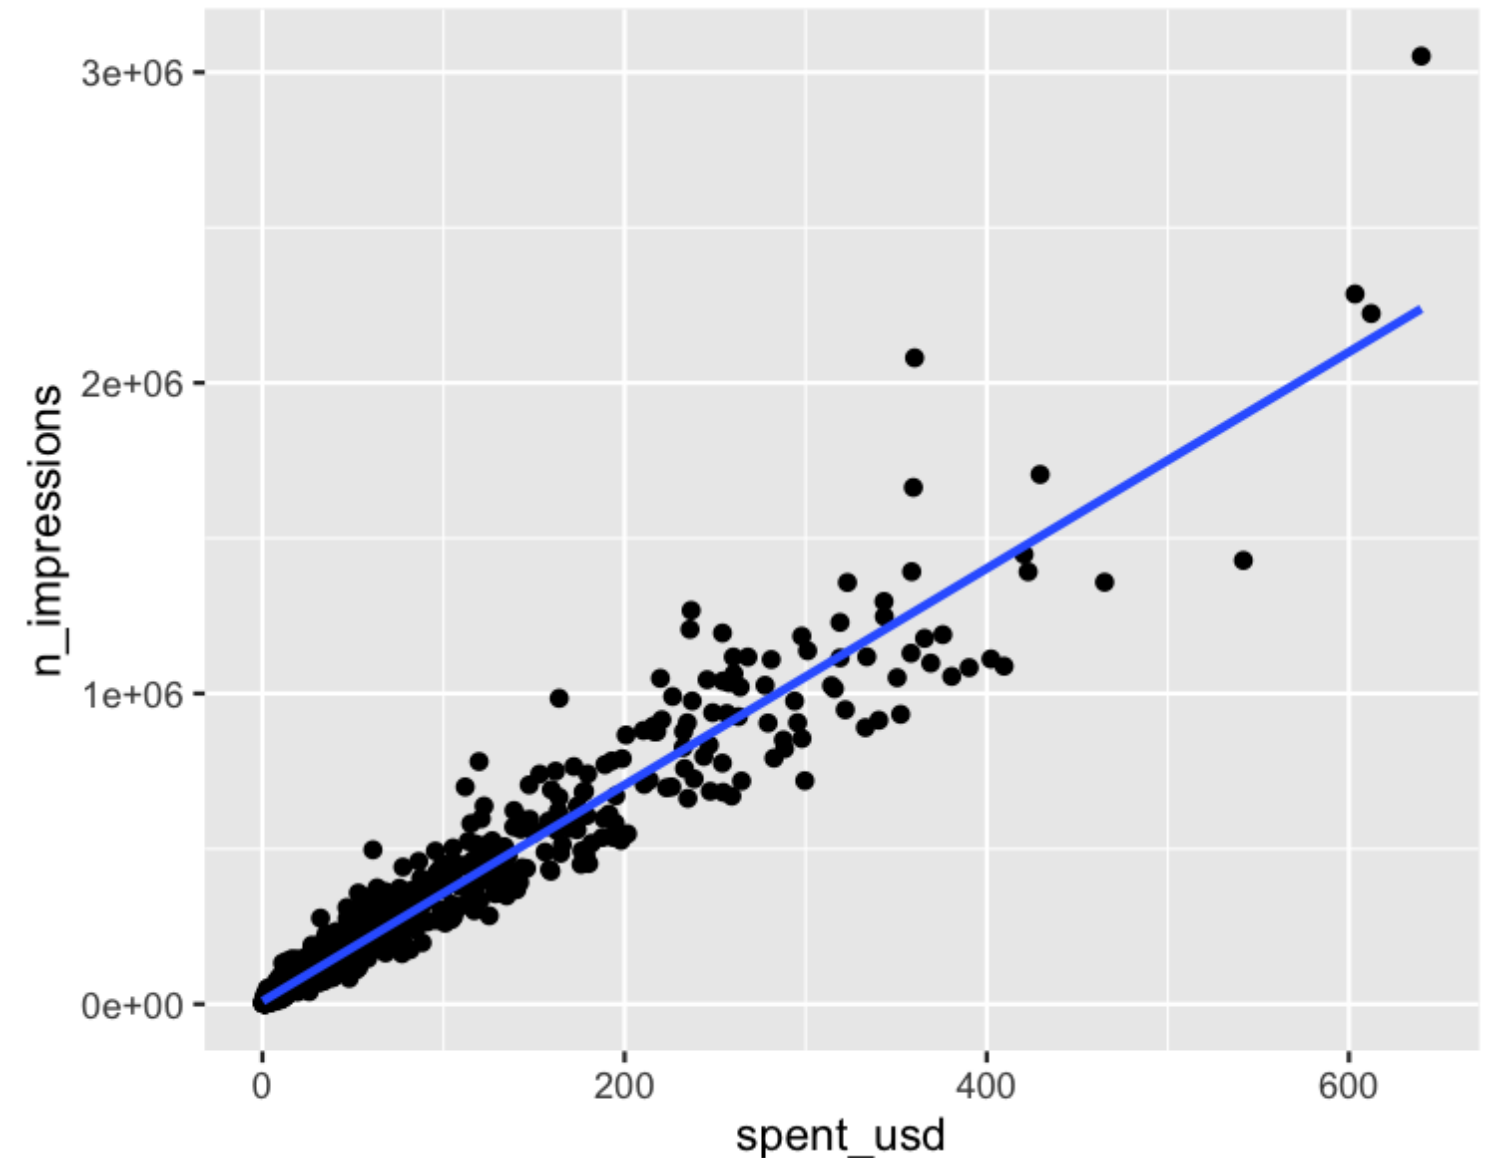
1. Pay Facebook to shows ads.
2. People see the ads ("impressions").
3. Some people who see it, click it.

- 936 rows
- Each row represents 1 advert

spent_usd	n_impressions	n_clicks
1.43	7350	1
1.82	17861	2
1.25	4259	1
1.29	4133	1
4.77	15615	3
...

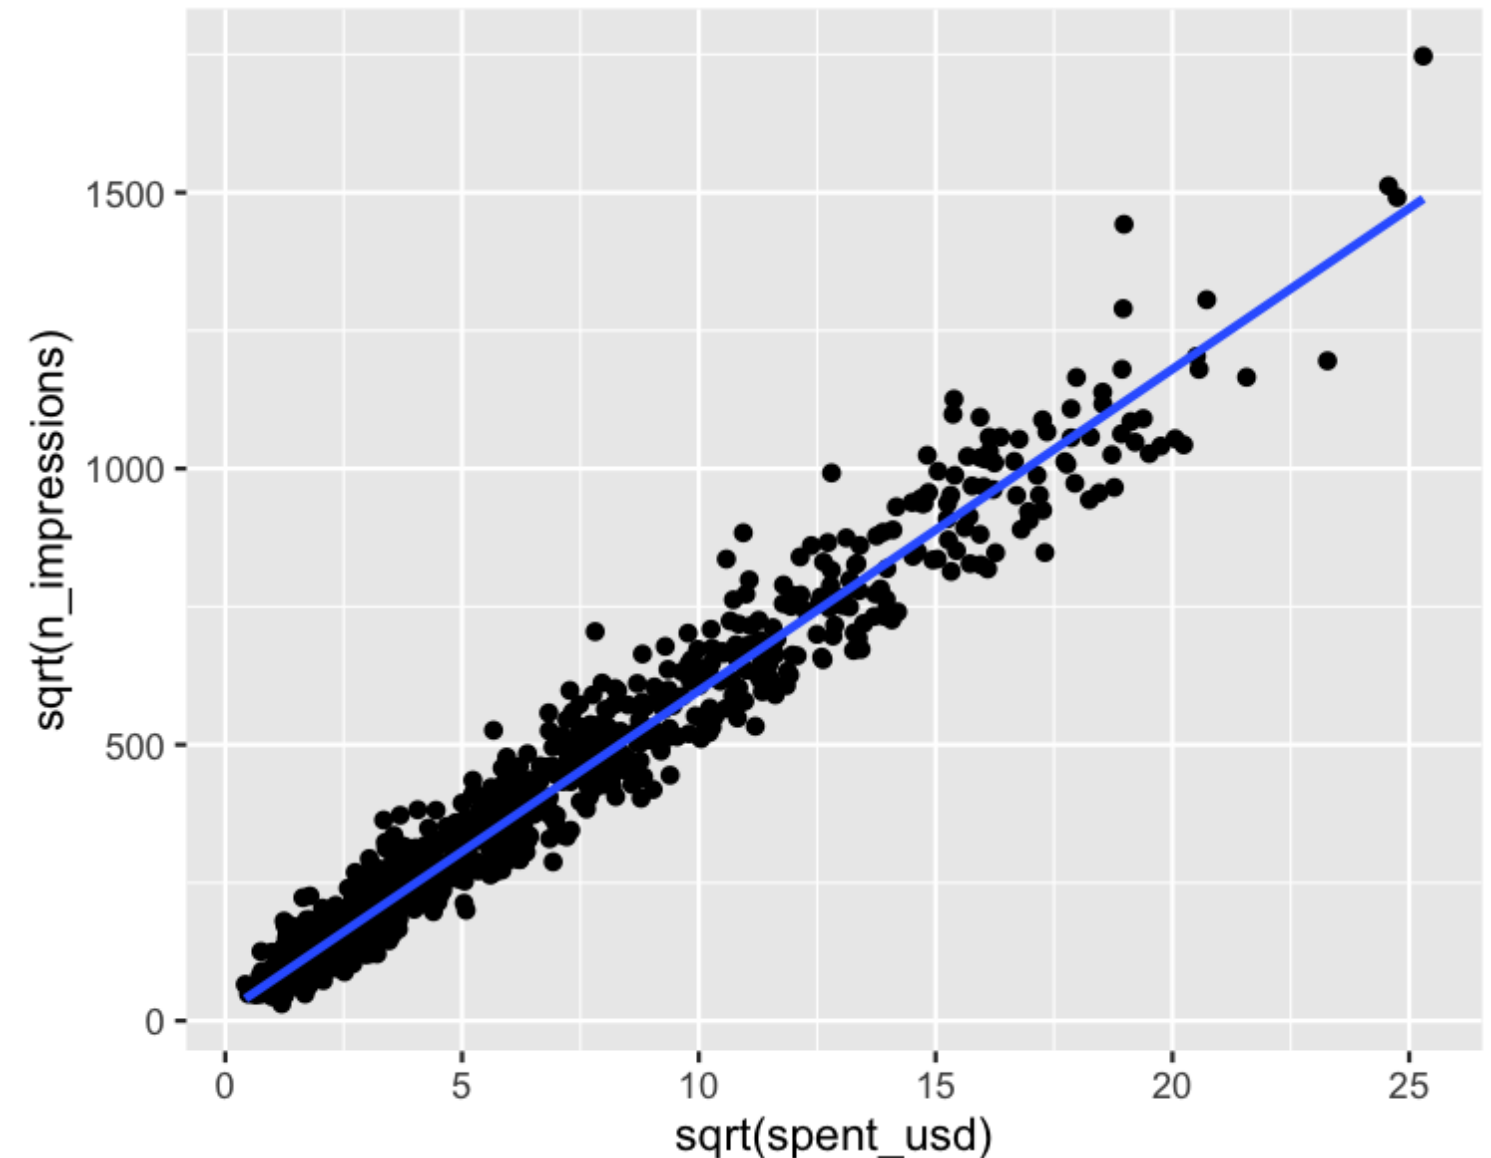
Plot is cramped

```
ggplot(  
  ad_conversion,  
  aes(spent_usd, n_impressions)  
) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE)
```



Square root vs square root

```
ggplot(  
  ad_conversion,  
  aes(sqrt(spent_usd), sqrt(n_impressions))  
) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE)
```



Modeling and predicting

```
mdl_ad <- lm(  
  sqrt(n_impressions) ~ sqrt(spent_usd),  
  data = ad_conversion  
)
```

```
explanatory_data <- tibble(  
  spent_usd = seq(0, 600, 100)  
)
```

```
prediction_data <- explanatory_data %>%  
  mutate(  
    sqrt_n_impressions = predict(  
      mdl_ad, explanatory_data  
    ),  
    n_impressions = sqrt_n_impressions ^ 2  
  )
```

```
# A tibble: 7 x 3  
  spent_usd sqrt_n_impressions n_impressions  
    <dbl>         <dbl>         <dbl>  
1         0          15.3           235.  
2        100         598.        357289.  
3        200         839.        703890.  
4        300        1024.       1048771.  
5        400        1180.       1392762.  
6        500        1318.       1736184.  
7        600        1442.       2079202.
```

Let's practice!

INTRODUCTION TO REGRESSION IN R